
CPSC 66 Final Report: Classifying Texts by Time Period

Zachary Kelly
Sasha Casada
Nathan Le

ZKELLY1@SWARTHMORE.EDU
SCASADA1@SWARTHMORE.EDU
NLE1@SWARTHMORE.EDU

Abstract

This project aims to classify historical texts by their period of writing, useful in digital humanity studies, such as history, anthropology, political science, and other fields, wherein dates of important documents remain unidentified. Traditional methods of classifying texts by time period require human experts with extensive knowledge of history and literature, which can be time-consuming and expensive. In our project, we investigated the effectiveness of the Naive Bayes, Random Forest, and ADABOOST models as accurate, precise, and efficient solutions to these problems. Using $n_{EXAMPLES} \approx 2000$ examples spanning $n_{CLASSES} = z$ classes, namely, $> 1700s, 1700-1799, 1800-1899, 1900-2000, > 2000$. We found all models to have some predictive capability, with *RandomForests* performing the best, *AdaBoost* performing below, but with potential for better performance, and *NaiveBayes* performing the worst among the three.

1. Introduction

This paper explores the effectiveness of three supervised ML models at classifying the time period of English written works. Given the broad potential application for this, it could prove to be useful in a number of fields. In developing our research, we primarily considered the field of Historical Dating; wherein historians sometimes find themselves unable to date an important document, or determine the author of a particular work from a few probable candidates due to the unavailability of a definite date of completion, etc. Accordingly we carefully designed our classification system with both of these problems in mind.

In order to accomplish this goal: We collected approx-

imately 2000 unique English texts dating from *pre-18th* century to *21st* century in the public domain from Project Gutenberg. We extracted features from this raw data relating to word-frequency, namely Bag-of-Words (BoW) and Short For-Term-Frequency Inverse Document Frequency (TF-IDF).

The three models we chose for this analysis are Naive Bayes, Random Forest, and ADABOOST. There has been extensive research into the application of Deep Learning to perform a variety of text-based tasks, but we wanted to explore the value use of these algorithms for this task due to their simplicity to implement, run, and test compared to Deep Learning Methods, especially with a relatively small and limited dataset.

2. Methods

2.1. Naive Bayes

The Naive Bayes model is a supervised model inherently based on Bayes' Theorem, that is, the probability of some event based on past information. It makes two fundamental assumptions, that is:

1. Each feature is independent of the others.
2. Each feature makes an equal contribution.

Because Naive Bayes makes the assumption that all features are independent makes the algorithm very fast compared to other algorithms for classification. For this reason in particular, it is most commonly applied to the task of classification of high-dimensional data (such as text classification). However, all features being independent of each other is not something that is often true in real life, thus the algorithm is often considered to be less accurate than other similar algorithms.

2.2. Random Forest

The Random Forest model is a supervised, ensemble model made up of a grouping of k decision trees that are created using randomly selected subsections of the training data.

A given novel example a is fed to the decision trees, their classifications are made, and the plurality decision is taken from their results.

There are a few properties of the Random Forest classifier that were desirable for our task:

1. Can handle high-dimensional data well.
2. Reduces overfitting; attributed to bagging and random feature selection.
3. It is rather robust against noise. Because text data can be noisy and contain irrelevant features, we would like an algorithm that can identify that most relevant features in the classification task.

2.3. AdaBoost

The AdaBoost model (short for: Adaptive Boosting) is a supervised, ensemble model made up of k decision stumps, that is, trees of depth 1. It attempts to build a stronger, more effective classifier from the k weaker classifiers. It does this by a cycle wherein it builds a model from the training data, then builds another model that improves upon the errors of the first one. This loop continues until all of the examples in the training set are predicted correctly or the loop reaches a max number of models.

Similar to above, there are also features specific to AdaBoost that we identified to be helpful in our task, key among them being its handling of class-imbalance. Present in our data-set are several imbalanced classes with labels that sometimes appear significantly more than others. Because of the minority class weight-assignment and use of weak classifiers, it can help us with those instances.

3. Experiments and Results

3.1. Data Collection

Requiring a large amount of English texts spanning a large period of time, we decided to use texts from the online Project Gutenberg Digital Archive. This site serves as an archive of digitized works currently in the public domain from a variety of fields, genres, and time periods. We utilized the open source (GPL-3.0) [Standardized Project Gutenberg Corpus \(SPGC\)](#) on GitHub in order to automatically scrape the text content and metadata associated with each UTF-8 work available on the site into a CSV. An additional script downloaded all of the works, provided word counts for each work, as well as tokenized versions.

3.2. Tokenization

For each work, we made use of the tokenized texts for our data processing. Tokenization is a crucial step for tasks

involving Natural Language Processing (NLP). It can be best described as the process for systematically breaking down a text into units (often called words, or subwords) which are used as inputs for various computational tasks, and in our case, feature extraction. To illustrate the concept, consider a text file which contains the following text:

“Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua...”

Tokenizing the above text and exporting this text to a file would get you the following:

```
1  lorem
2  ipsum
3  dolor
4  sit
5  amet
6  consectetur
7  ...
8
```

Listing 1. Tokenization Example

Now with individual words on each newline, as well as getting rid of any unnecessary whitespace, symbols, capitalization, etc. we become more easily able to process the content of the text.

3.3. Data Processing

Once we collected and labeled sufficient data for our project, we took a number of steps to process the texts. Firstly, we removed all non-English texts. As our primary goal was to predict the time period of English text, we removed all works translated to English from another language more than a decade past their original publication date so as to prevent these texts from interfering with our classifiers’ learning processes and/or affecting their validation results.

Next, we fine tuned our data to exclude duplicate texts, anthropologies and compilations of texts published over a period of time. The former to ensure these duplicates did not affect the learning of our classifiers or artificially inflate our dataset, and the latter was because of the nature of these texts, their contents would not necessarily accurately reflect their publication year.

When we arrived at an English-only, duplicate-free corpus, we then manually labeled all of our examples by using Google Scholar to find the ground-truth publication dates of the texts. We believed this necessary because, although Project Gutenberg kept detailed information on the publication dates of all entries, these sometimes did not align with the true publication date of the text. For instance, although *The Magna Carta* has been republished several times, thus there were several entries for it in the original dataset, none

of which with publication dates that were its correct publication date.

Then, we assembled approximately 100 texts into a single document and tokenized it. This was in order to produce a feature vector of standardized order by representing each bucket of the vector as a word ordered by its frequency. This would allow us to then feed other texts into the same feature generation algorithms and compare them with each other, as the features of a given text would be in the same order as the others. The implementation for this can be seen at `code/generate_wordFrequency.py` and is exported to `code/data`.

3.4. Feature Generation

We implemented two word vectorization algorithms that would assign a numeric feature value for a given word w in the vector. The vector is constructed such that words occurring with a higher frequency in the subsample text occur earlier in the vector.

3.4.1. BAG OF WORDS (BoW)

Bag-of-Words (Bow) is a numerical classification model where the frequency of occurrence of each word is used as a feature for training a classifier. In this case, each feature represents the total number of occurrences of a particular word in a text, in this case the subsample text.

To illustrate Bag of Words (BoW), consider the figure below which models the following two sentences:

Ex. 1 “The dog sat, the cat sat too.”

Ex. 2 “The dog and the cat and the other cat all sat.”

Text #	f_0 ‘the’	f_1 ‘dog’	f_2 ‘cat’	f_3 ‘sat’
Ex. 1	2	1	1	1
Ex. 2	3	1	2	1

3.4.2. TERM-FREQUENCY INVERSE DOCUMENT FREQUENCY (TF-IDF)

Is a numerical statistic that is intended to reflect how important a word is to a document in a collection or corpus. TF-IDF value increases proportionally to the number of times a word appears in the document and is offset by the number of documents in the corpus that contain the word.

This can be expressed mathematically as:

$$W_{i,j} = tf_{i,j} \times \log \frac{(N)}{df_i}$$

3.5. Data Generation, Train, and Test Pipeline

We used 1850 example texts, with a total of 20,000 features. We vectorized these using TF-IDF, which came out

to be the more effective (in terms of accuracy), vectorization algorithm as compared to Bag of Words. For labels, we went and used century-long periods for our slicing, and we removed stopwords from our vectorization output. Stopwords are usually very common words that appear frequently in a language and do not add much meaning to the text, such as “the”, “a”, “an”, “and”, etc.

Once we generated our data, we ran our dataset through our three algorithms, as described above, with validation.

3.6. Validation

3.6.1. NAIVE BAYES CLASSIFIER

Hyperparameter tuning was performed on the α parameter across the range (0, 1] This parameter space was explored with Randomized Search with 10 sampling iterations and 3 fold cross validation.

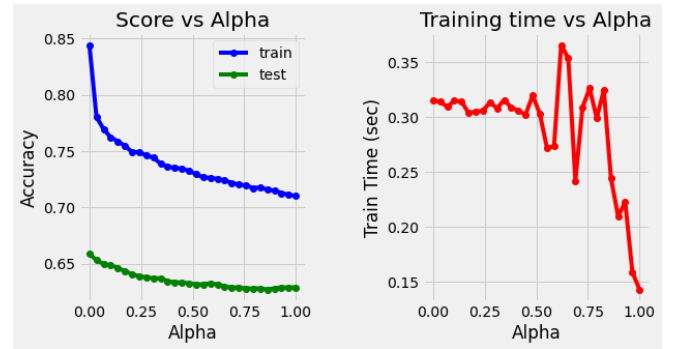


Figure 1. Validation curve of α for Naive Bayes Classifier

Table 1. tuned parameters for naive bayes

parameter	value
alpha	0.011633196056020756

3.6.2. RANDOM FOREST CLASSIFIER

Hyperparameter tuning was performed on the following parameters:

- $n_estimators$: the number of trees in the forest, with a range from 200 to $\frac{NumFeatures}{3}$.
- $max_features$: the number of features to consider when looking for the best split, with options ‘auto’ and ‘sqrt’.
- max_depth : the maximum number of levels in the tree, ranging from 10 to 110.
- $min_samples_split$: the minimum number of samples required to split a node, with a range of values [2, 10].

- *min_samples_leaf*: the minimum number of samples required at each leaf node, with a range of values [1, 4].
- *bootstrap*: the method of selecting samples for training each tree, with options of either True or False.

This hyperparameter space was explored with Random Search with 10 sampling iterations and 3 fold cross validation.




Figure 2. Validation curve of *max_depth* for Random Forest Classifier




Figure 3. Validation curve of *n_estimators* for Random Forest Classifier

Table 2. tuned parameters for Random Forest Classifier

parameter	value
<i>n_estimators</i>	918
<i>min_samples_split</i>	10
<i>min_samples_leaf</i>	1
<i>max_features</i>	sqrt
<i>max_depth</i>	60
<i>bootstrap</i>	False

3.6.3. ADABOOST CLASSIFIER

Hyperparameter tuning was performed on the following parameters for the AdaBoost classifier:

- *n_estimators*: the number of weak learners to train iteratively, with a range from 50 to 500.
- *learning_rate*: the learning rate, which determines the contribution of each weak learner, ranging from 0.001 to 1.0.
- *base_estimator*: the type of base estimator used as a weak learner, in this case, a decision tree classifier with varying maximum depths from 1 to 15.

This parameter space was explored with Random Search with 10 sampling iterations and 3 fold cross validation.




Figure 4. Validation curve of *n_estimators* for AdaBoost Classifier

Table 3. Tuned Parameters for AdaBoost Classifier

Parameter	Value
<i>n_estimators</i>	250
<i>learning_rate</i>	0.9388367346938776
<i>base_estimator</i>	<i>max_depth</i> = 7

3.7. Results

Accuracy was used as the primary performance metric for our three models. After parameter tuning, the performance changes are as follows:

Table 4. Model Performance

Model	Base	Tuned	Performance Delta
AdaBoost	0.5189	0.7270	0.2081
Naive Bayes	0.6527	0.7054	0.0527
Random Forest	0.7568	0.7773	0.0205

The Naive Bayes and Random Forest models experienced rapidly diminishing returns on increasing model complexity for higher accuracy. It is possible that AdaBoost can be fitted with a greater number of estimators for higher accuracy, if not for the high computational runtime.






Figure 6. Confusion Matrix for Random Forest Classifier

The confusion matrices for the models are as follows:

All three models show that their predictions heavily bias towards either the 1800s or the 1900s. This not surprising given the high class imbalance of our dataset.

The confusion matrices for the Naive Bayes and Random Forest classifiers imply these models are splitting their predictions predominantly between texts $< 1700s - 1899$ and $1900 - > 2000s$ with decent accuracy. This implies that although these models can capture differences between these two time periods, the models struggle at making higher resolution classifications.

The AdaBoost confusion matrix shows that, though its predictions bias towards 1800s or 1900s, it is able to somewhat accurately classify texts outside of the 1800 and 1900

range. Notably, it does a decent job at classifying texts < 1700 and $1700 - 1799$. This implies that, despite the class imbalance, the AdaBoost classifier is able to model some time specific patterns of older texts.

We can inspect the patterns these trees are modeling by examining the feature importance of the Random Forest and AdaBoost classifiers. The following graphs represent the Gini Impurity / Mean Decrease in Impurity for the features of the classifiers.




Figure 8. Feature Importance Random Forest

The most notable difference between these two graphs is that the Random Forest classifier has a more even distribution of MDI across the features while the AdaBoost classifier has a sharper decrease in MDI across the features. Interestingly, seemingly innocuous words (e.g. poor and later) have the highest MDI values. In addition to more ar-




Figure 9. Feature Importance AdaBoost

chaic words (e.g. hath and alas) having significant MDI values. This could imply that the decision trees initially split on commonly found words, and deeper into the trees do they diverge into more obviously time specific language.

4. Social Implications

Our classification system has the potential to greatly impact the digital humanities community. Once unsolved problems could now, with some degree of reliability, be solved. Furthermore, as our system is not proprietary, those most interested in the subject may freely and openly use and improve it. Therefore, it provides a resource to the digital humanities community that allows the individual to make discoveries, not just, say, a research group or well-funded corporation, providing an opportunity for those with deep interest in the subject to contribute when they may not have been able to before.

That being said, a research group or corporation itself is made up of some number of individuals passionate about the material, and if this classification system is significantly better than their performance, this could lead to lower funding for these individuals.

5. Challenges

We faced a number of challenges in designing and implementing our classification system. They can be broken down into the following categories:

5.1. Design and Data Collection

While our first challenge of finding a sufficient quantity of suitable data was very easily solved, even the solution proved to have some difficulty. Due of the nature of the Project Gutenberg archive's structuring of meta-data asso-

ciated with texts, it was not possible for us to be able to find dates of original publish for each of the works that we encountered. As a solution, we resorted to manual labeling of data, which was high cost in terms of time-investment.

5.2. Implementation

Implementing the models also came with challenge. The first and most substantial was creating some standardized feature vector that allowed us to compare texts to each other. Without this, we would not be able to compare documents at all, as the same bucket in the feature vector for one document could stand for another feature in another document. After deliberation on how to best solve this problem, we developed the aforementioned standardized vector from a compilation of texts.

5.3. Testing and Validation

In the case of training a model to do classification of texts based on something as abstracted to the natural development of language as time period is, you end up with a problem where you are unable to accurately predict texts that should otherwise be classifiable because of the boundaries between two target labels being too close to a particular text's date of writing. An example would be that of a two texts, one written in 1899, and one written in 1900. Because of the cut-off boundary arbitrarily being between centuries, two texts which likely have more in common with each other than other texts of the same class as them end up being classified as two very different things. Coming up with solutions on how to mitigate this boundary error could help improve with accuracy over predictions.

6. Conclusions

Despite the high class imbalance of our dataset, thee relatively simply models have shown promising results in text classification by time period. The Naive Bayes and Random forest models do well at distinguishing between texts ;1899 and ;1900, but are largely unable to classify texts outside of those time periods. The AdaBoost classifier seems to be able to classify texts outside of this time range, and it is likely that with a larger and better balanced dataset, and higher complexity (e.g more estimators) that its performance may increase.

6.1. Class Imbalance Issues

Our original idea was to have several labels for a range of time periods, in order to better perfect our identification methods. However, we have a very bad case of sample balancing. This sample balancing is not something that could have been fixed by undersampling (this would cut down an extreme amount over the data that we have available), or

oversampling (this would lead to extreme overfitting over the data that we do have available since we do not have a sufficient number of examples for these less sampled periods). Accordingly, we have cut down on the range of dates that we are processing. Instead, we are opting to do analysis per half-century between 1800s and 1950s.

7. Future Directions

We believe that there is lots of room for improvement, but we have still accomplished good work in the way of our initial goal. There still exists issues, related to the project that we have encountered that existed outside of the scope of what our plan developed into. However, if we were to continue development outside of our four weeks, this is what we would consider:

7.1. Feature Embedding

Since we have such a high number of features, having some sort of feature embedding/feature encoder feature reduction system would be very good for actually reducing the amount of data that we need to process. Because of the most common words appearing within a certain number of texts, there are many books which will simply exclude.

7.2. Sample Rebalancing

As mentioned in the conclusions section, we still have quite a few issues related to class balance. Perhaps we might be able to find higher accuracy over some of the more difficult to estimate time periods by sampling more evenly in terms of our data.

7.3. Deep Learning

With a larger and better balanced dataset, it may be possible to apply deep learning techniques to model the time differences in texts. Deep learning models are well equipped to capture spatio-temporal patterns in data, which may allow them to perform better at this type of classification

7.4. Feature Semantic Analysis

The feature importance of the Random Forest and AdaBoost classifiers are measured by their Gini Index. For a feature space of high cardinality, this may not be the best representation of feature importance. Using different measures of feature importance, performing semantic analysis on important features may yield interesting results on what words or sets of words are most relevant at classifying texts (e.g. the set of words relating to modern technology, or a set of old timey customary words)

Acknowledgments

We would like to acknowledge our gratitude to the extensive and helpful oversight, advice, and support of Prof. Ben Mitchell to this project, as well as the Swarthmore Computer Science Department for granting us access to department resources for this project.

References

Langley, P. Crafting papers on machine learning. In Langley, Pat (ed.), *Proceedings of the 17th International Conference on Machine Learning (ICML 2000)*, pp. 1207–1216, Stanford, CA, 2000. Morgan Kaufmann.